

# Geohash in Golang Assembly

Lessons in Absurd Optimization

30 August 2018

Michael McLoughlin

Software Engineer, Uber ATG



# What is Geohash?

# Geohash is a Point Encoding Scheme

Maps a latitude longitude pair to a 64-bit integer or equivalent string.

- Indexing geo data
- Proximity queries
- Used in Redis, for example

Example:

- **Gophercon** is at: (39.74279, -104.99706)
- Integer geohash is: 0xeb7f254240fd612
- String geohash is: [tuvz4p141zc1](#)

Geohash can be taken at various *precisions*.

Divides the earth into *grid cells* of various sizes.

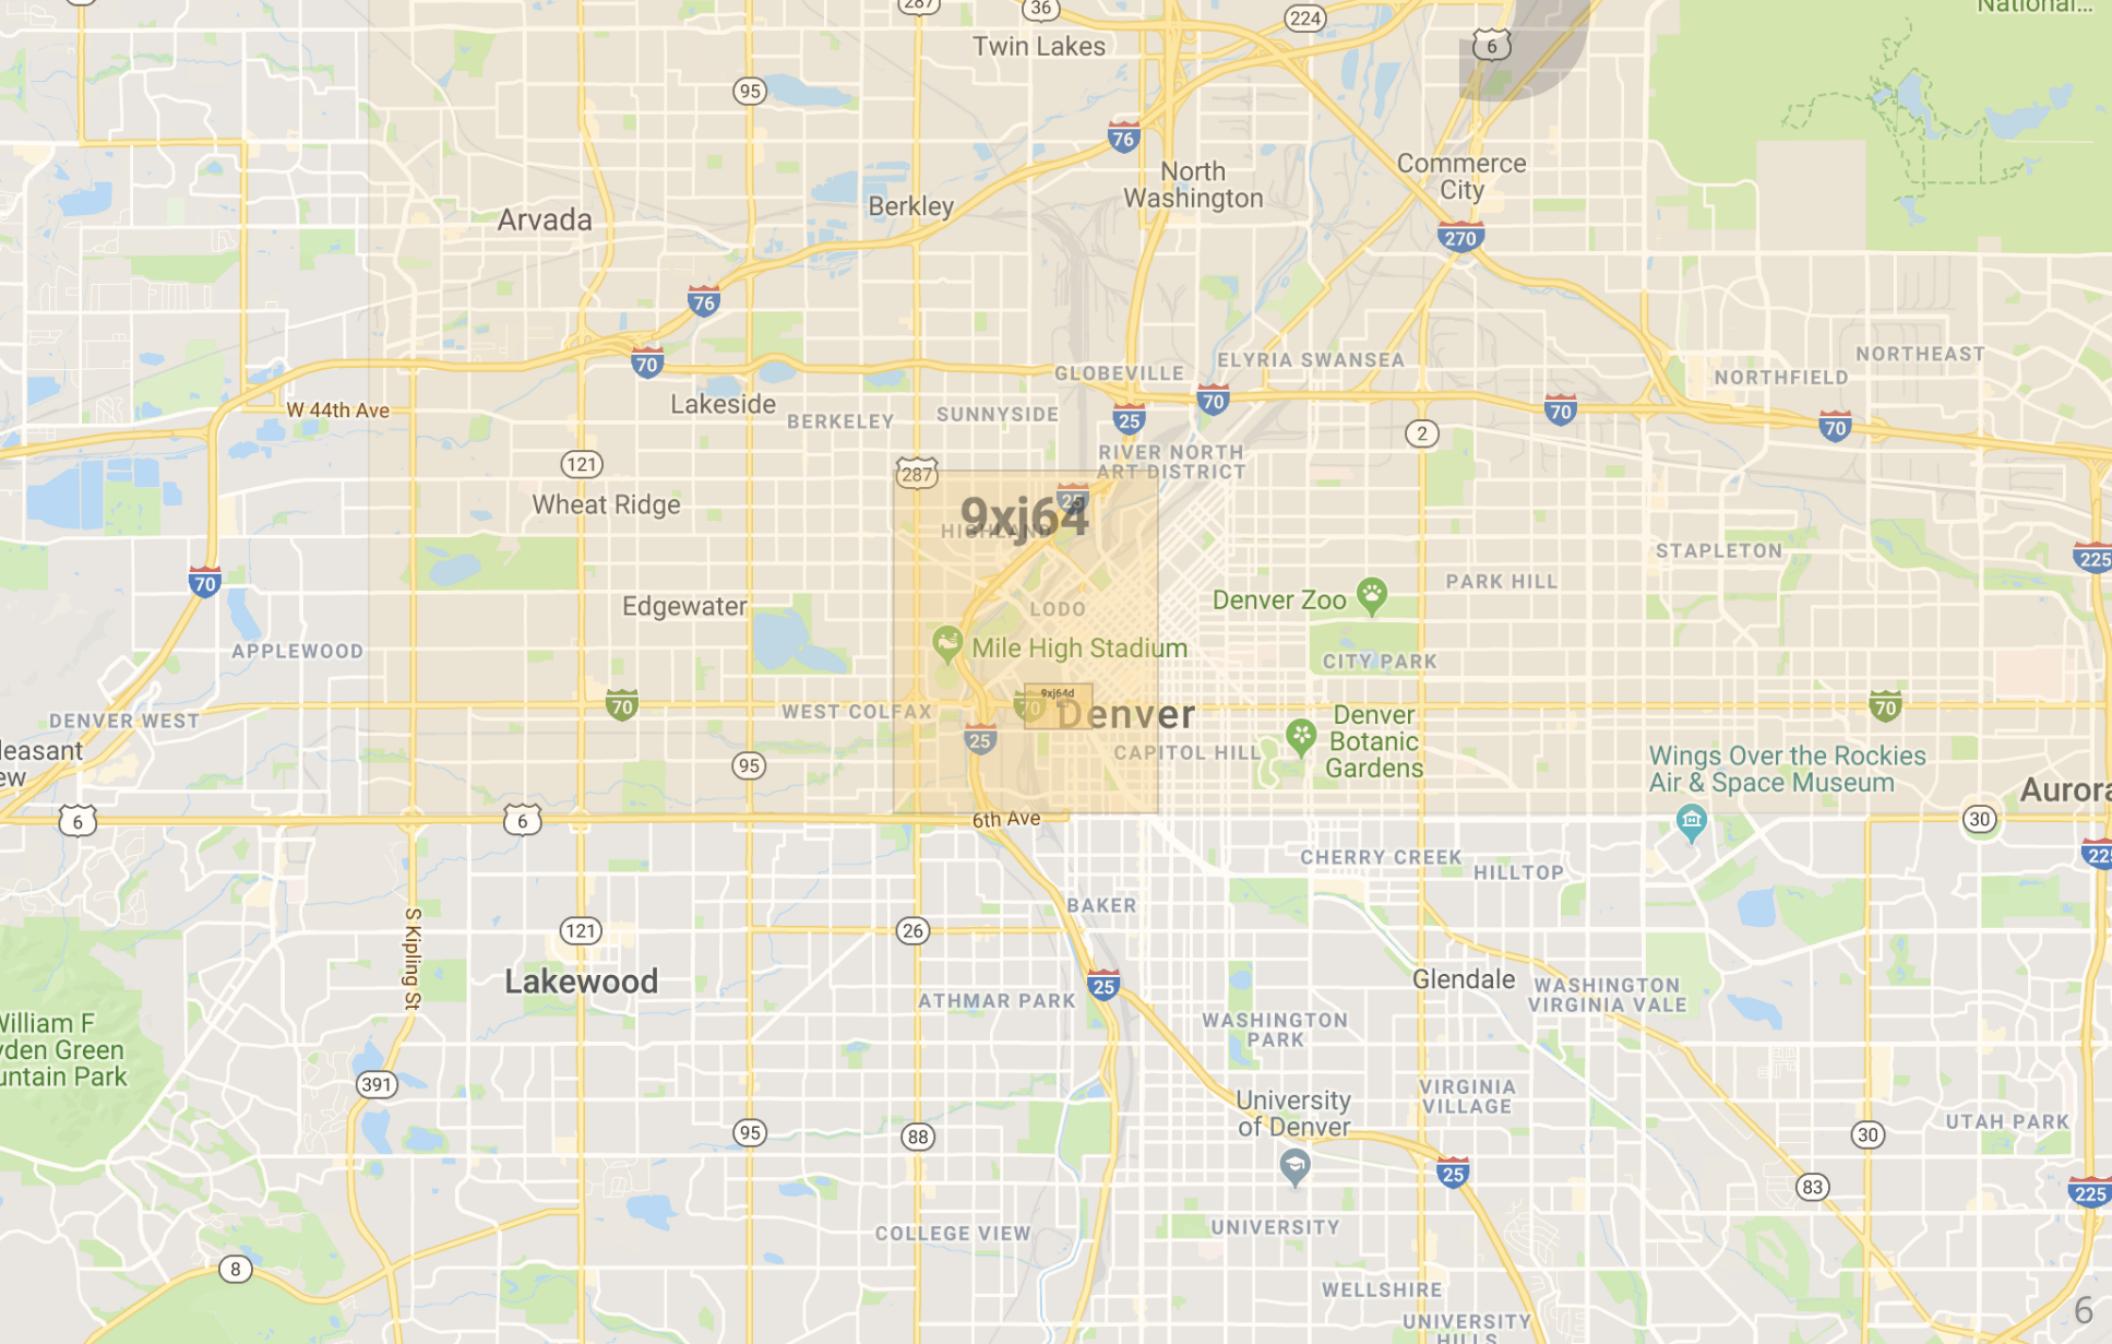
9xi64dz

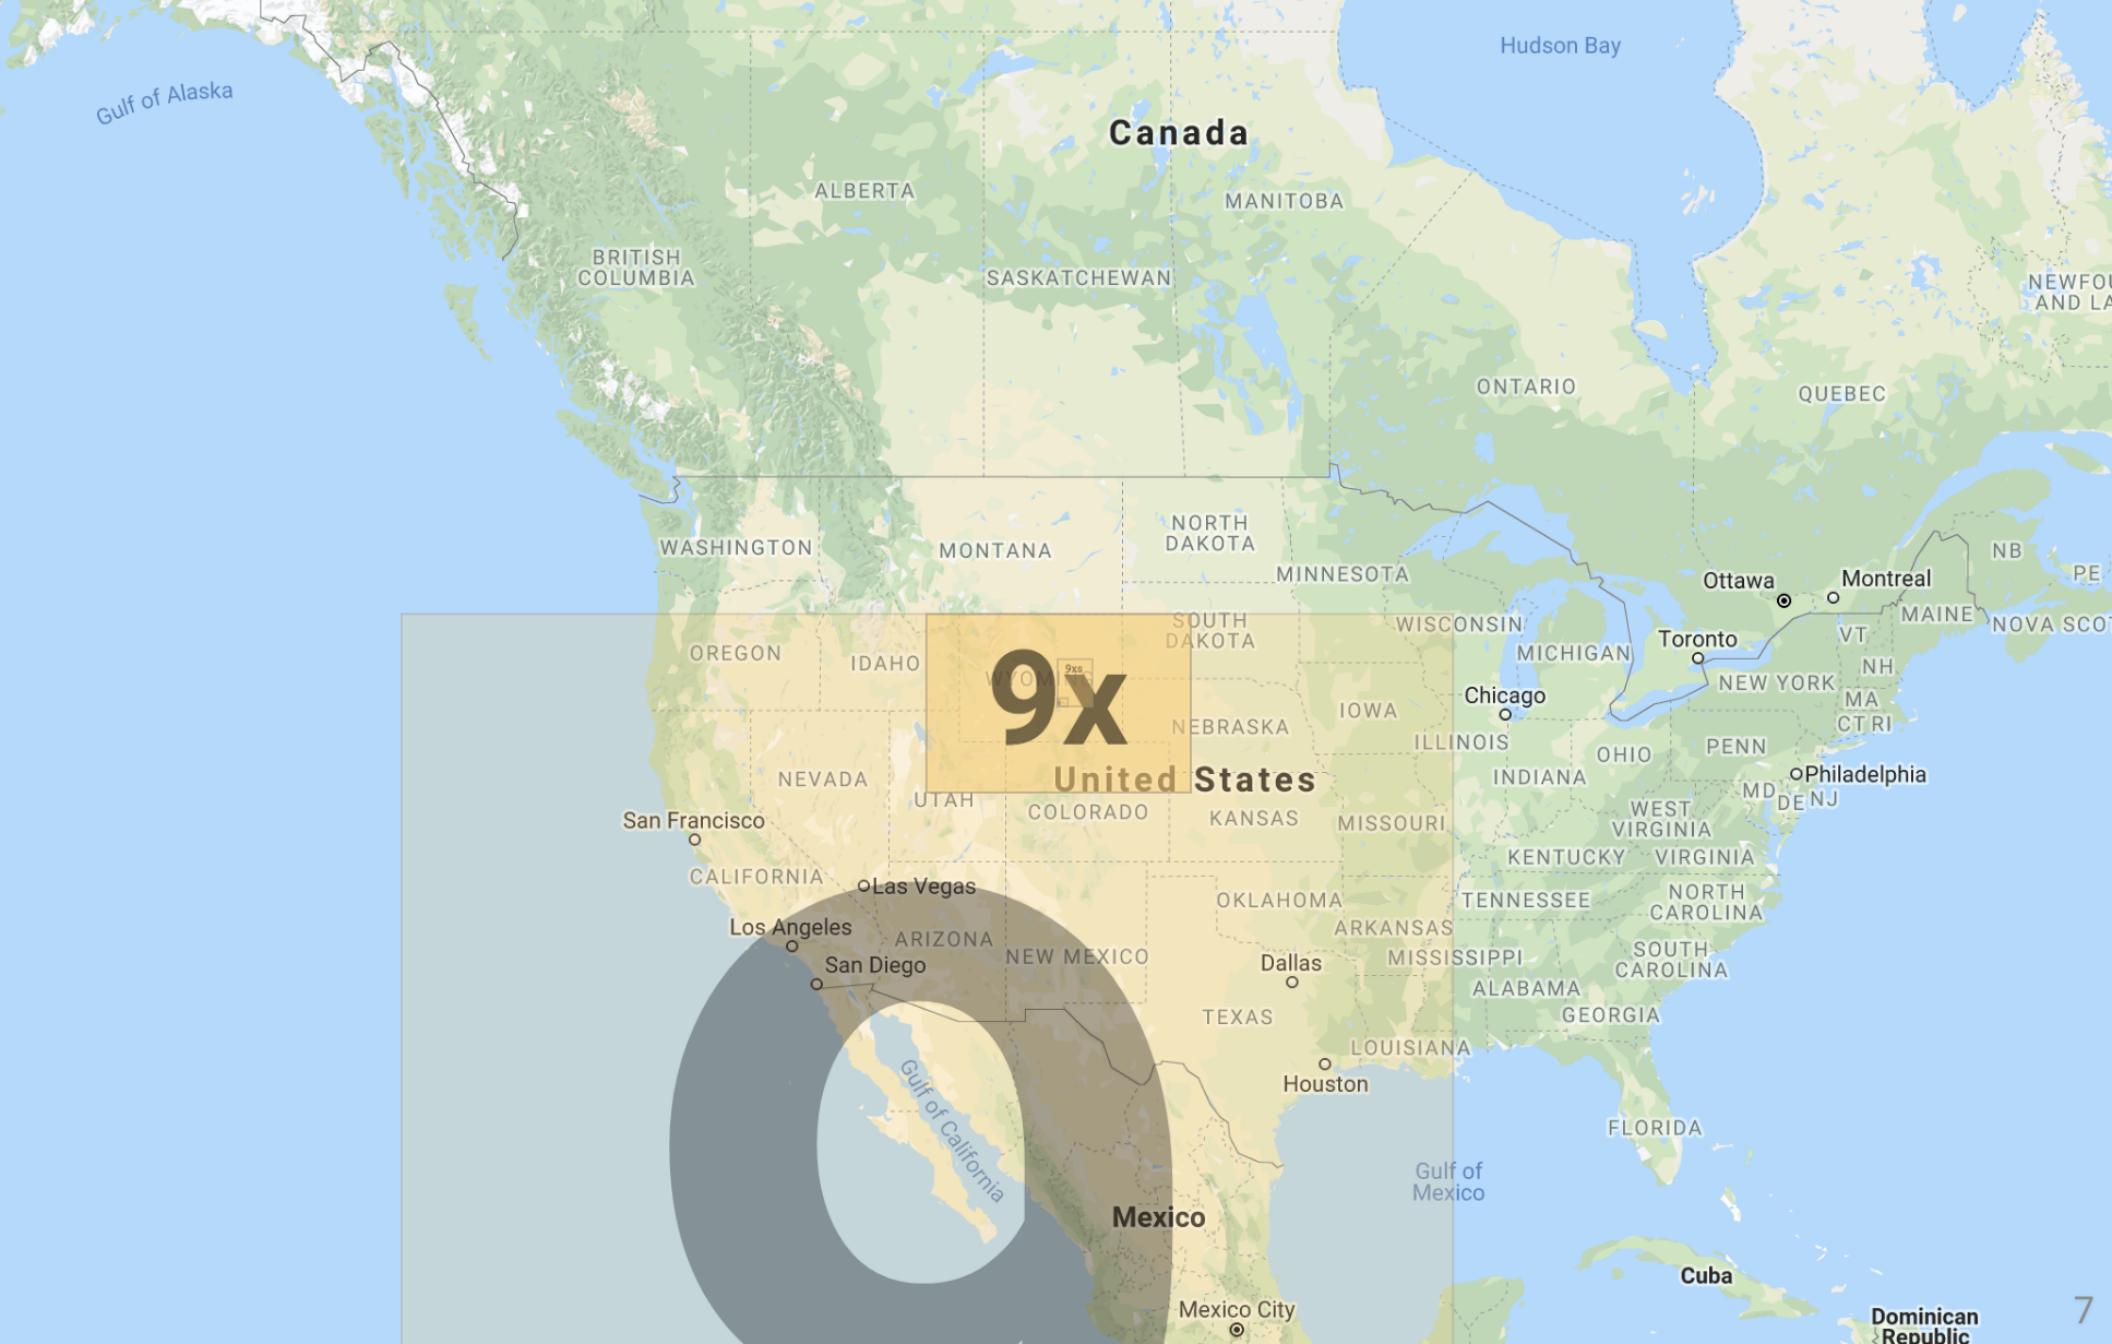
9xj64dzq

Mile High Ballroom  
Theater

The UPS Store

National...







# Forget about String Geohashes

For *internal* or performance-sensitive use cases integer form is enough.

# How it Works

## Step 1: Quantize latitude/longitude to 32-bits

39.74279

-104.99706

1	0	1	0	0	1	1	1	1	0	0	1	1	1	0	0	0	1	1	1	1	0	0	1	0	0
a	7	c	e	2	3	e	4																		

1	0	1	1	1	1	1	0	1	1	1	0	1	0	0	0	0	0	1	0	0	0	1	1	1	0	0
b	d	d	d	0	4	3	9	1																		

- Input: (39.74279, -104.99706)
- Quantized: (0xa7ce23e4, 0xbdd04391)

## Step 2: Bit Interleave

39.74279

-104.99706

1	0	1	0	0	1	1	1	1	0	0	1	1	1	0	0	0	1	0	0	1	0	0
a		7		c		e		2		3		e		4								

1	0	1	1	1	1	1	0	1	1	1	0	1	0	0	0	0	0	1	0	0	0	1
b		d		d		d		0		4		3		9		1						

1	1	0	0	1	1	1	0	1	0	1	1	1	1	1	0	0	1	0	1	0	0	0	1
c		e		b		7		f		2		5		4		2		4		0		f	d

- Input: (39.74279, -104.99706)
- Quantized: (0xa7ce23e4, 0xbdd04391)
- Interleaved: 0xeb7f254240fd612

This is the *integer geohash*.

# Pure Go

## High Level

```
// EncodeInt encodes the point (lat, lng) to a 64-bit integer geohash.  
func EncodeInt(lat, lng float64) uint64 {  
    return Interleave(Quantize(lat, lng))  
}
```

# Quantize

39.74279

-104.99706

1	0	1	0	0	1	1	1	1	0	0	0	1	0	0	0	1	0	0
a	7	c	e	2	3	e	4											

1	0	1	1	1	1	0	1	1	1	0	1	0	0	0	0	0	1	0	0
b	d	d	d	0	4	3	9	1											

```
// Quantize maps latitude and longitude to 32-bit integers.  
func Quantize(lat, lng float64) (lat32 uint32, lng32 uint32) {  
    lat32 = uint32(math.Ldexp((lat+90.0)/180.0, 32))  
    lng32 = uint32(math.Ldexp((lng+180.0)/360.0, 32))  
    return  
}
```

## Intermediate Step: Spread

39.74279

-104.99706

1	0	1	0	0	1	1	1	1	0	0	1	1	1	0	0	0	1	0	0	1	1	1	1	0	0	1	0	0
a	7	c	e	2	3	e	4																					

1	0	1	1	1	1	0	1	1	1	0	1	0	0	0	0	1	0	0	0	1	1	1	1	0	0	1	0	0
b	d	d	d	0	4	3	9	1																				

0	1	0	0	0	1	0	0	0	0	1	0	1	0	1	0	1	0	0	0	0	0	0	0	0	0	0	0
0	1	0	0	0	1	0	1	0	1	0	1	0	1	0	1	0	0	0	0	0	0	0	0	0	0	0	0

1	1	0	0	1	1	1	0	1	1	0	1	1	1	1	1	1	0	0	1	0	1	0	1	1	1	1	0	1	0	1
c	e	b	7	f	2	5	4	2	4	0	f	d	6	1	2															

```
// Spread out the 32 bits of x into 64 bits, where the bits of x occupy even
// bit positions.
func Spread(x uint32) uint64 {
    X := uint64(x)
    X = (X | (X << 16)) & 0x0000ffff0000ffff
    X = (X | (X << 8)) & 0x00ff00ff00ff00ff
    X = (X | (X << 4)) & 0x0f0f0f0f0f0f0f0f
    X = (X | (X << 2)) & 0x3333333333333333
    X = (X | (X << 1)) & 0x5555555555555555
    return X
}
```

# Interleave

39.74279

1	0	1	0	0	1	1	1	0	0	1	1	1	0	0	0	1	0	0	1	1	1	1	0	0	1	0	0
a	7	c	e	2	3	e	4																				

-104.99706

1	0	1	1	1	0	1	1	1	0	1	0	0	0	0	1	0	0	0	1	1	1	1	0	0	1	0	0
b	d	d	0	4	3	9	1																				

0	1	0	0	0	1	0	0	0	0	1	0	1	0	1	0	1	0	0	0	0	0	0	1	0	0	0
0	1	0	0	0	1	0	1	0	1	0	1	0	1	0	1	0	0	0	0	0	0	0	1	0	0	0

1	1	0	0	1	1	1	0	1	0	1	1	1	1	1	1	0	0	1	0	0	0	0	1	1	1	1
c	e	b	7	f	2	5	4	2	4	0	f	d	6	1	2											

```
// Interleave the bits of x and y. In the result, x and y occupy even and odd  
// bitlevels, respectively.  
func Interleave(x, y uint32) uint64 {  
    return Spread(x) | (Spread(y) << 1)  
}
```

27.40

Nanoseconds / Geohash

36.50

Million Geohashes / Second

## Pure Go Performance

# Assembly

## 17 Instructions

```
TEXT ·EncodeIntAsm(SB), NOSPLIT, $0
    MOVSD lat+0(FP), X0
    MOVSD lng+8(FP), X1

    MULSD $(0.00555555555555556), X0
    ADDSD $(1.5), X0
    MOVQ X1, R9
    SHRQ $20, R9

    MULSD $(0.002777777777777778), X1
    ADDSD $(1.5), X1
    MOVQ X0, R8
    SHRQ $20, R8

    MOVQ $0x5555555555555555, BX
    PDEPQ BX, R8, R11
    PDEPQ BX, R9, R10
    SHLQ $1, R10
    XORQ R10, R11

    MOVQ R11, ret+16(FP)
    RET
```

## How to draw an owl

1.



2.



1. Draw some circles

2. Draw the rest of the  
owl

# Scale

```
TEXT ·EncodeIntAsm(SB), NOSPLIT, $0
```

```
    MOVSD lat+0(FP), X0
```

```
    MOVSD lng+8(FP), X1
```

```
MULSD $(0.00555555555555556), X0
```

```
    ADDSD $(1.5), X0
```

```
    MOVQ X1, R9
```

```
    SHRQ $20, R9
```

```
MULSD $(0.00277777777777778), X1
```

```
    ADDSD $(1.5), X1
```

```
    MOVQ X0, R8
```

```
    SHRQ $20, R8
```

```
    MOVQ $0x5555555555555555, BX
```

```
    PDEPQ BX, R8, R11
```

```
    PDEPQ BX, R9, R10
```

```
    SHLQ $1, R10
```

```
    XORQ R10, R11
```

```
    MOVQ R11, ret+16(FP)
```

```
    RET
```

# Unit Interval Trick

```
TEXT ·EncodeIntAsm(SB), NOSPLIT, $0
```

```
    MOVSD lat+0(FP), X0
```

```
    MOVSD lng+8(FP), X1
```

```
    MULSD $(0.00555555555555556), X0
```

```
    ADDSD $(1.5), X0
```

```
    MOVQ X1, R9
```

```
    SHRQ $20, R9
```

```
    MULSD $(0.002777777777777778), X1
```

```
    ADDSD $(1.5), X1
```

```
    MOVQ X0, R8
```

```
    SHRQ $20, R8
```

```
    MOVQ $0x5555555555555555, BX
```

```
    PDEPQ BX, R8, R11
```

```
    PDEPQ BX, R9, R10
```

```
    SHLQ $1, R10
```

```
    XORQ R10, R11
```

```
    MOVQ R11, ret+16(FP)
```

```
    RET
```

# PDEP

```
TEXT ·EncodeIntAsm(SB), NOSPLIT, $0
    MOVSD lat+0(FP), X0
    MOVSD lng+8(FP), X1

    MULSD $(0.00555555555555556), X0
    ADDSD $(1.5), X0
    MOVQ X1, R9
    SHRQ $20, R9

    MULSD $(0.002777777777777778), X1
    ADDSD $(1.5), X1
    MOVQ X0, R8
    SHRQ $20, R8

    MOVQ $0x5555555555555555, BX
    PDEPQ BX, R8, R11
    PDEPQ BX, R9, R10
    SHLQ $1, R10
    XORQ R10, R11

    MOVQ R11, ret+16(FP)
    RET
```

# Interleave

```
TEXT ·EncodeIntAsm(SB), NOSPLIT, $0
```

```
    MOVSD lat+0(FP), X0
```

```
    MOVSD lng+8(FP), X1
```

```
    MULSD $(0.00555555555555556), X0
```

```
    ADDSD $(1.5), X0
```

```
    MOVQ X1, R9
```

```
    SHRQ $20, R9
```

```
    MULSD $(0.002777777777777778), X1
```

```
    ADDSD $(1.5), X1
```

```
    MOVQ X0, R8
```

```
    SHRQ $20, R8
```

```
    MOVQ $0x5555555555555555, BX
```

```
    PDEPQ BX, R8, R11
```

```
    PDEPQ BX, R9, R10
```

```
    SHLQ $1, R10
```

```
    XORQ R10, R11
```

```
    MOVQ R11, ret+16(FP)
```

```
    RET
```

2 . 32

Nanoseconds / Geohash

431 . 03

Million Geohashes / Second

## Assembly Implementation

# SIMD

- What if we use special instructions to do four geohashes at once
- Intel Intrinsics
- Interleaving trick thanks to [Daniel Lemire](#)

```
TEXT ·EncodeIntSimd(SB), NOSPLIT, $0
    MOVQ lat+0(FP), AX
    MOVQ lng+24(FP), BX
    MOVQ hash+48(FP), CX

    VBROADCASTSD reciprocal180<>+0x00(SB), Y0
    VMULPD      (AX), Y0, Y0
    VBROADCASTSD onepointfive<>+0x00(SB), Y1
    VADDPD      Y1, Y0, Y0
    VPSRLQ      $20, Y0, Y0
    VBROADCASTSD reciprocal360<>+0x00(SB), Y2
    VMULPD      (BX), Y2, Y2
    VADDPD      Y1, Y2, Y1
    VPSRLQ      $20, Y1, Y1
    VMOVDQU     spreadbyte<>+0x00(SB), Y2
    VPSHUFB     Y2, Y0, Y0
    VBROADCASTSD lonibblemask<>+0x00(SB), Y3
    VPAND       Y3, Y0, Y4
    VMOVDQU     spreadnibblelut<>+0x00(SB), Y5
```

1 . 13

Nanoseconds / Geohash

883 . 00

Million Geohashes / Second

## SIMD Implementation

## Conclusion

- Massive performance gains are possible
- Generally prefer to rely on the compiler
- Assembly may make sense to access of special instruction sets

Full details:

- [Geohash in Golang Assembly](#) blog post
- [mmccloughlin/geohash](#) package
- [Comparison of all Golang geohash packages](#)

Thank you

Michael McLoughlin  
Software Engineer, Uber ATG  
[mmcloughlin@gmail.com](mailto:mmcloughlin@gmail.com)  
<https://mmcloughlin.com/>  
[@mbmcloughlin](https://twitter.com/mbmcloughlin)